

<p>Massive data processing Homework 1 Gaétan Morand</p>

Configuration

Mac OS 10.11

Hadoop 2.7.3 installed with Homebrew

The standard output does not indicate the total CPU time, so I have to report the times with the UNIX « time » command.

a.i

To collect stopwords we use the same map function as a simple Wordcount Mapreduce program. The difference is that we use the reducer to filter only the words which appear more than 4,000 times.

In order to use 10 reducers we must add the line:

```
job.setNumReduceTasks(10);
```

Due to that the results are divided into 10 individual files, which we shall merge with the following command:

```
hdfs dfs -getmerge output ouput.csv
```

The program reports an execution time of 36.816s.

Here are the identified stopwords :

a	at	day	good	how	let	most	of	s	t	thing	us	who
about	away	did	got	i	like	mr	old	said	take	think	very	will
after	be	do	great	if	little	much	on	say	than	this	was	with
again	been	down	had	in	ll	must	one	see	that	thou	way	would
all	before	ever	has	into	long	my	only	shall	the	thy	we	yet
am	but	every	have	is	made	never	or	she	their	time	well	you
an	by	first	he	it	make	no	other	should	them	to	were	your
and	can	for	her	its	man	not	our	sir	then	too	what	
any	come	from	here	just	may	nothing	out	so	there	two	when	
are	could	get	him	king	me	now	over	some	these	up	where	
as	d	go	his	know	more	o	own	such	they	upon	which	

They are also in the file a.i.csv.

The code is in the file a.i.java.

a.ii

We add a combiner that does the same job as the reducer, except for the filtering part. The combiner adds up the counts of each word, but we leave it to the reducer to only keep the most used. Indeed this operation must be executed only at the end of the process.

```
public static class Combine extends Reducer<Text, IntWritable, Text, IntWritable> {  
    public void combine(Text key, Iterable<IntWritable> values, Context context)  
        throws IOException, InterruptedException {  
        int sum = 0;  
        for (IntWritable val : values) {  
            sum = sum + val.get();  
        }  
        context.write(key, new IntWritable(sum));  
    }  
}
```

The resulting csv file is the same as the precedent one, but the program reports an execution time of 35.912s.

It is slightly faster this way because the combiner is called right after the map function, which is highly parallelized. In reality my processor only has 2 cores so parallelization does not add much.

The code is in the file a.ii.java.

a.iii

In order to compress the map results, I added the following lines:

```
conf.set("mapreduce.map.output.compress", "true");  
conf.set("mapreduce.map.output.compress.codec",  
"org.apache.hadoop.io.compress.SnappyCodec");
```

The program reports an execution time of 36.837s. It is longer than before. Indeed the limiting factor on my computer is the processing power, not the transfer time between memory and processor. So we waste more time compressing the data than we save transporting it.

The code is in the file a.iii.java.

a.iv

To use 50 reducers, we must add this line:

```
job.setNumReduceTasks(50);
```

The program reports an execution time of 40.850s. This is because my computer is only able to run 2 threads at a time since it has 2 cores. So there is no point in having more than 2 threads. It does not compute faster with 50 threads than with 10. It is even slower since we waste computing power dividing and merging the data, and managing these threads.

b.

To learn how to skip word, I read this website:

https://www.cloudera.com/documentation/other/tutorial/CDH5/topics/ht_wordcount3.html

There are two major components to build in this program: the part that skips stopwords in the map function, and the part that remove the duplicate filenames, in the reduce function.

The output file is in the file b.txt

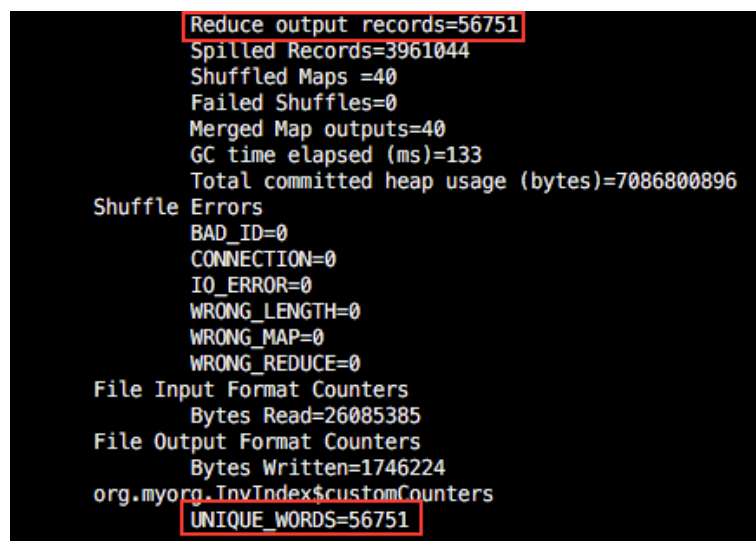
The code is in the file b.java

c.

The total number of words is indicated by the counter “Reduce output records”. In this case there are 56,751 different words.

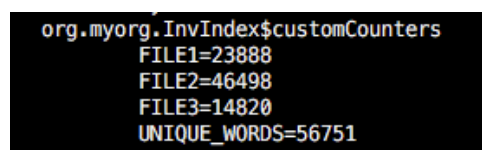
I created a new counter, which increments in each instance of the reduce method. It confirms that there are 56,751 different words.

Here is a screenshot of the output:



```
Reduce output records=56751
Spilled Records=3961044
Shuffled Maps =40
Failed Shuffles=0
Merged Map outputs=40
GC time elapsed (ms)=133
Total committed heap usage (bytes)=7086800896
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
Bytes Read=26085385
File Output Format Counters
Bytes Written=1746224
org.myorg.InvIndex$customCounters
UNIQUE_WORDS=56751
```

Then we have to create a new counter for each file. We can only count in the Reduce method because we want to avoid duplicates. We test for the name of the file.



```
org.myorg.InvIndex$customCounters
FILE1=23888
FILE2=46498
FILE3=14820
UNIQUE_WORDS=56751
```

The code is in the file c.java

The output file is counters.txt

d.

Here we have to take back the code from our inverted index. The difficulty is to maintain a different wordcount for each file. To do this I chose to use a Map class, in which the key represents the file name, and the value represents the number of occurrences.

I am sure it can also be done using two mapreduces sequentially but I am not sure how to program this.

I had a bug in the final problem which I uploaded.

```
java.lang.Exception: java.lang.ClassCastException: org.apache.hadoop.io.Text cannot be cast to java.util.HashMap
    at org.apache.hadoop.mapred.LocalJobRunner$Job.runTasks(LocalJobRunner.java:462)
    at org.apache.hadoop.mapred.LocalJobRunner$Job.run(LocalJobRunner.java:529)
Caused by: java.lang.ClassCastException: org.apache.hadoop.io.Text cannot be cast to java.util.HashMap
    at org.myorg.InvIndex$Reduce.reduce(InvIndex.java:114)
    at org.myorg.InvIndex$Reduce.reduce(InvIndex.java:95)
    at org.apache.hadoop.mapreduce.Reducer.run(Reducer.java:171)
    at org.apache.hadoop.mapred.ReduceTask.runNewReducer(ReduceTask.java:627)
    at org.apache.hadoop.mapred.ReduceTask.run(ReduceTask.java:389)
    at org.apache.hadoop.mapred.LocalJobRunner$Job$ReduceTaskRunnable.run(LocalJobRunner.java:319)
    at java.util.concurrent.Executors$RunnableAdapter.call(Executors.java:511)
    at java.util.concurrent.FutureTask.run(FutureTask.java:266)
    at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1142)
    at java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:617)
    at java.lang.Thread.run(Thread.java:745)
```

I was unfortunately not able to debug this.

I uploaded the code in the file d.java. I think it is really close to working, but I have to upload now to meet the deadline.